

Ayushman Srivastava

+1 (412)-450-7558 | Email: mailto:ayushman@gmail.com | LinkedIn: www.linkedin.com/in/ayushman28/ | GitHub: github.com/ay-2814/

EDUCATION

Carnegie Mellon University (Graduated, August 2024)

Master of Science in information Systems Management, *Machine Learning & Artificial Intelligence Track*

Pittsburgh, PA

August 2024

Courses: Deep Learning, Machine Learning in Production, Computational Data Science, A/B Testing, Applied Machine Learning, Distributed Systems, Data Structures and Algorithms, Object Oriented Programming in Java, Probability and Statistics

SKILLS

Machine Learning and OOP: Python, PyTorch, TensorFlow, SciPy, nltk, langchain, Scikit-learn, NumPy, Pandas, R, Java, C++, C#, Java

Distributed and Cloud Computing: CUDA, MapReduce, Spark, Hadoop, AWS, GCP, Azure, Kafka, Docker, Kubernetes, Airflow

Scaling ML Systems: VertexAI, Django, flask, MLFlow, Prometheus, Grafana, WandB, DVC, Cron, Jenkins, GitHub

Database: SQL, MySQL, PostgreSQL, Snowflake, MongoDB, NoSQL **Visualization:** PowerBI, Tableau, Excel VBA, React Native

WORK EXPERIENCE

Season Labs – Co-founder and Chief Technical Officer

October 2025– Present

- CTO of a venture backed startup that is making an empirical model for user intent when interacting with AI agents

CaseFlood.ai – Co-founder and Head of AI (Y Combinator Winter 2025)

September 2024– September 2025

- Co-founded a legal AI voice startup that was backed by Y combinator and was part of its Winter 2025 batch
- Built the product from ground up and was pivotal in attracting a \$80 million Series A offer and before that, securing \$3.2 million seed round for the company at a \$30M valuation

Swiggy (India's largest Food and grocery delivery app) — Data Science Manager, Trust & Safety Charter

June 2022 - May 2023

- Orchestrated experiments and applied statistical inferences to detect new signals and features for **fraud attribution**
- Built **ensemble models** based on XGBoost, Multilinear regression etc. utilizing **40+ real-time signals** to detect new fraud types, performed **multiple ad-hoc analysis**, resulting in detection of 5 new fraud types and savings of **\$12.5 million** in first 6 months
- Developed Fraud Scorecard, by analyzing order level data of ~30 million users, and 500K daily orders, derived multiple metrics to place customers into suitable segments, recommended strategies to improve **Customer Satisfaction (CSAT)** score by 0.4 points
- Led weekly Business Reviews with VPs and C-suite, creating executive summaries and actionable insights to improve KPIs

Mu Sigma Inc

September 2018 - May 2022

Chief Data Scientist for Computer Vision app for a Fortune 100 CPG client

- Technical lead of 10 developers and ML Engineers to build iOS app for Pest & Disease Management of Cocoa Plants, managed efforts for **data annotation** using Scale.ai APIs based on taxonomy, built **Image recognition** and data collection functionalities.
- Applying **YOLO** and **Canny Edge detection** on drone images to detect and classify the plants and pods with change in color
- Developed the app for 500+ agricultural scientists to collect, analyze, and visualize data for a 400-hectare farm, saving 25 hours/week per person in manual data collection and improving pest detection by ~40%.

Senior Data Scientist working with a leading CPG client

- Oversaw team of 8 Data Engineers and Analysts to deliver an end-to-end sales reporting solution for 15+ markets
- Reduced time** for analyzing sales force data and generating recommendations by **12 hours per week per user**

Data Scientist for a leading Consumer Packaged Goods (CPG) client

- Led a team of 4 data analysts to optimize a portfolio of **1200+** Stock Keeping Units (SKU) across 8+ markets globally
- Developed a **Prescriptive model**, utilizing **Hypothesis testing** and **Statistical inference**, delivering **\$7.4 million in profits**, and paved the way for the solution to be scaled further

RELEVANT PROJECTS

Attention-based speech recognition leveraging Transformer architecture

March 2024 – May 2024

- Developed encoder and decoder modules including positional encoding (from the *Attention Is All You Need* paper), multi-head attention and feed-forward projections, combined them all to construct the full transformer model
- Experimented with different configurations of model hyperparameters, such as the number of layers, model dimensionality, and the number of attention heads, achieving a **Levenshtein distance of 16**

Automated Movie Recommender System Using Collaborative Filtering

January 2024 - May 2024

- Created and maintained end-to-end pipeline for an automated movie recommender system in production with event stream data (Kafka) of **1 million users** and more than **30,000 movies**.
- Configured model to recommend 20 movies for each request with max **latency of 500ms**, maintaining uptime of **99.99%**
- Conducted **A/B testing** and canary testing between 2 models in production and reported the differences with 95% confidence

LEADERSHIP

- Impact Award at Mu Sigma for leading 10 membered team deliver successful projects with constant scope changes 2022
- Vice-President, Heinz Analytics Club CMU 2023-2024
- Runner-up: Google Hackathon on Sustainability Goals CMU 2024